

Dhananjay G.

MSc Computer Science, University of Glasgow | ETL/ELT Pipeline Design | 4 years of experience | Data Engineer London, UK | +44 7393568854 | guptadhananjay090@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Google Cloud Certified Data Engineer with expertise in building **cloud-native architecture** across **Azure**. Known for excellence in **ETL**, **Python**, **SQL**, and **Pyspark**, delivering impact such as a **£1M+ revenue uplift**, and **40% faster support resolution** using GenAI tools. Led a **migration of 50M+ records**, improving data cleaning from **80% to 99.9%**. Received special letters of recommendation from industry leaders and academics for excellence in managing data.

RELEVANT SKILLS

- **Functional Competencies:** Data Modelling, Data Pipeline Development, Data Quality Improvement, Data Governance & Security Compliance, Mentoring & Coaching, Cross-Functional Collaboration, Stakeholder Engagement, Data Culture Advocacy, Continuous Learning
- **Technical Tools:**
 - **Programming:** Python (pandas, Numpy), SQL, PySpark, Bash
 - **Data Engineering Frameworks:** Apache Spark, SparkSQL, Airflow
 - **Cloud Platforms:** Azure (ADLS Gen2, Data Factory, Databricks, Analytics), GCP (BigQuery, Cloud Storage)
 - **GenAI Tools & Frameworks:** OpenAI, Gemini, LangGraph, RAG Architecture, Tokenisation
 - **Development & DevOps:** Git, GitLab, Docker, JIRA, Azure DevOps, CI/CD, Github
 - **Databases:** MS SQL (T-SQL), Vector database(Qdrant, Pinecone), PostGRES, NoSQL

WORK EXPERIENCE

Data Engineer, Agilitics

Jan 2025 - Present

- **Built and maintained an automated ELT pipeline for a client using Azure Data Factory and SQL Database**, ingesting over 3M+ finance records from an on-premises Oracle system into a cloud-based Azure data platform, supporting incremental loads (medallion architecture) in the staging layer.
- Engineered transformation and change-data capture (CDC) logic using T-SQL and MD5 hashing within Azure SQL and **Delta Lake (parquet format) on Azure Databricks**, ensuring high data fidelity across revenue and billing datasets, and enabling accurate, audit-ready finance reporting at scale.
- **Delivered a production-grade finance data mart and Power BI reporting solution** by modeling segmented customer- and company code-level metrics, powering dashboards for active customer counts, average rate per unit, and month-over-month general ledger trends, supporting financial planning and analysis across 20+ regional business units.
- **Designed and deployed a scalable retrieval-augmented generation (RAG) chatbot (FastAPI)** by leveraging **LangChain, OpenAI, and Qdrant**, enabling over 500 users to query internal documents (unstructured data) in natural language, **reducing support ticket resolution time by 40%**.

Data Engineer, Open Insight

Apr 2022 - Jun 2023

- **Designed and implemented 40+ business-critical KPIs** by analysing **400K user records for 2nd largest telecoms in Thailand** using **PySpark and Azure data lake storage(ADLS Gen2)** on **Azure Databricks**, directly informing targeting strategies that **generated £500K in incremental revenue**.
- **Boosted pipeline performance by 50%** by redesigning **data processing logic**, optimising joins, and applying strategic **caching** within Azure Databricks, reducing overall runtime from **35+ hours to under 12**.
- **Built automated, production-grade ETL workflows** that ingested data from **15+ complex sources**, applying **SQL (T-SQL)** and **Delta Lake transformations** to ensure scalability, data consistency, and **60% faster processing speeds**.
- **Developed end-to-end PySpark pipelines** to process over **400K telecom customer records**, enabling real-time KPI tracking, web (socials) and app customer insight, and **marketing dashboard refreshes via Power BI** with significantly improved timeliness.
- **Integrated KPI pipelines into GitLab CI workflows**, enabling continuous deployment and automated table refreshes in production; also **mentored two interns**, accelerating their upskilling and helping ensure consistent,

on-time delivery across sprints.

Associate Analyst, Mastech Infotrellis

Sep 2020 - Mar 2022

- **Improved data accuracy from 80% to 99.9%** by leading the end-to-end migration for DELL technologies of **50M+ records** (customer, address, credit) from Oracle databases (structured data) to a unified global system across EMEA, APJ, and LATAM (data warehouse), using robust validation and reconciliation methods.
- **Reduced batch processing time by 50%** (4hrs to 2hrs) by automating ETL pipelines in Python and SQL(T-SQL) with REST API integration, flag-based incremental logic, and parallelised regional workflows that scaled efficiently across large datasets.
- **Minimised post-migration issues and ensured schema compliance** by designing transformation workflows with clearly defined business rules, data cleaning checks, and unit testing integrated into GitLab CI/CD using pytest.
- **Resolved ad-hoc data issues and unblocked critical processes** by collaborating closely with developers, testers, and analysts, proactively addressing tasks like custom customer data seeding and migration bottleneck triage.
- **Accelerated environment rollout by 30%** by managing sprint deliverables in Azure DevOps within an Agile setup, aligning technical teams and stakeholders collaboration through structured planning, task tracking, and cross-functional coordination while creating technical documentation.

PROJECTS

- **LLM-Based Data Quality Analyser** (GenAI, FastAPI, OpenAI, Hugging Face, Streamlit, AWS bedrock): Built a GenAI tool to analyse CSV and Parquet files stored in Amazon S3, detecting nulls, outliers, and structural anomalies using FastAPI and Pandas. Enabled **70% faster QA checks** by generating natural language summaries and automated fixes through a user-friendly Streamlit UI, with exportable reports for business teams.
- **Customer360** (Azure Data Factory, Azure Databricks, ADLS Gen2, PySpark, SQL, Key Vault) Built an event-triggered pipeline to validate and route incoming CSV order files using ADF and Databricks, securely managing credentials via Azure Key Vault. Automated the ingestion and transformation of multi-source datasets using PySpark and SQL, generating **clean customer-level outputs** that supported advanced analytics and decision-making.

EDUCATION

MSc. Computer Science, University of Glasgow

Sep 2023 - Sep 2024

Highlight: Developed scalable data pipelines and backend systems as part of academic projects, including a cloud-based ETL solution using Python, SQL, and Azure that processed 5M+ records, supporting real-time analytics use cases.

B.Tech. Computer Science, SRM University

Aug 2016 - Aug 2020

Highlight: Built a mini data warehouse using MySQL and Python for a student performance tracking system, streamlining grade analysis and reducing report generation time by 60% for academic staff.

CERTIFICATIONS & RECOGNITIONS

- **Google Cloud Professional Data Engineer**, Issued by Google Cloud, July 2025
- **Database and SQL for Data Science**, Issued by IBM, 2024
- **Data Analysis with Python**, Issued by IBM, 2024
- **GitLab Certified Associate**, Issued by GitLab, 2024
- Special Letter of Recommendation Received from Dr. Megha, Professor, SRM University
- Special Letter of Recommendation Received from Dr. Tarun Yadav, Researcher, DRDO (Gov. of India)
- Special LinkedIn Recommendation Received from Louis Sebastian, Principal SW Engineer, DELL Technologies